

main

December 3, 2020

```
[18]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb
import matplotlib.patches as mpatches
import collections
from collections import Counter
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import KFold, StratifiedKFold
from sklearn.preprocessing import RobustScaler
from sklearn.model_selection import StratifiedShuffleSplit
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import ShuffleSplit
from sklearn.model_selection import learning_curve
from sklearn.pipeline import make_pipeline
from imblearn.pipeline import make_pipeline as imbalanced_make_pipeline
from imblearn.over_sampling import SMOTE
from imblearn.under_sampling import NearMiss
from sklearn.metrics import roc_curve
from sklearn.model_selection import cross_val_predict
from sklearn.metrics import roc_auc_score
from sklearn.metrics import recall_score
from sklearn.metrics import precision_score
from sklearn.metrics import f1_score
from sklearn.metrics import accuracy_score
from sklearn.metrics import average_precision_score
import pickle
from sklearn.metrics import mean_squared_error
import warnings
warnings.filterwarnings("ignore")
```

```

[22]: data = pd.read_csv('creditcard.csv')
rob = RobustScaler()
data['rs_amount'] = rob.fit_transform(data['Amount'].values.reshape(-1,1))
data['rs_time'] = rob.fit_transform(data['Time'].values.reshape(-1,1))
data.drop(['Time','Amount'],axis=1, inplace=True)

rs_amount = data['rs_amount']
rs_time = data['rs_time']

data.drop(['rs_amount', 'rs_time'],axis=1, inplace=True)
data.insert(0, 'rs_amount', rs_amount)
data.insert(1, 'rs_time', rs_time)
x = data.drop('Class', axis=1)
y = data['Class']
skf = StratifiedKFold(n_splits=5, random_state=None, shuffle=False)

for tr_index, te_index in skf.split(x, y):
    orig_x_train, orig_x_test = x.iloc[tr_index], x.iloc[te_index]
    orig_y_train, orig_y_test = y.iloc[tr_index], y.iloc[te_index]
orig_x_train = orig_x_train.values
orig_y_train = orig_y_train.values
orig_x_test = orig_x_test.values
orig_y_test = orig_y_test.values

data= data.sample(frac=1)
F_data = data.loc[data['Class'] == 1]
NF_data = data.loc[data['Class']==0][:492]
sub_data = pd.concat([F_data,NF_data]).sample(frac = 1, random_state = 42)
sub_y = sub_data['Class']
sub_x = sub_data.drop('Class', axis = 1)
train_x, test_x, train_y, test_y = train_test_split(sub_x, sub_y, test_size=0.
    ↪2, random_state=42)
train_x = train_x.values
train_y = train_y.values
test_x = test_x.values
test_y = test_y.values
us_x = data.drop('Class', axis=1)
us_y = data['Class']
# the following make pipelined codes was modified from www.kaggle.com/
↪janiobachmann
for tr_index, te_index in skf.split(us_x, us_y):
    us_xtrain, us_xtest = us_x.iloc[tr_index], us_x.iloc[te_index]
    us_ytrain, us_ytest = us_y.iloc[tr_index], us_y.iloc[te_index]

us_xtrain = us_xtrain.values
us_xtest = us_xtest.values
us_ytrain = us_ytrain.values

```

```


us_ytest = us_ytest.values

Pkl_Filename_LR_U = "Pickle_LR_undersample.pkl"
Pkl_Filename_LR_O = "Pickle_LR_oversample.pkl"
Pkl_Filename_KNN = "Pickle_KNN_undersample.pkl"
Pkl_Filename_DTC = "Pickle_DTC_undersample.pkl"
with open(Pkl_Filename_LR_U, 'rb') as file:
    LOR_clf = pickle.load(file)
with open(Pkl_Filename_LR_O, 'rb') as file:
    best_OS_LOR = pickle.load(file)
with open(Pkl_Filename_KNN, 'rb') as file:
    KNN_clf = pickle.load(file)
with open(Pkl_Filename_DTC, 'rb') as file:
    decision_tree_clf= pickle.load(file)
# the following make pipelined codes was modified from www.kaggle.com/
→janiobachmann
for train, test in skf.split(us_xtrain, us_ytrain):
    us_pipeline =
    →imbalanced_make_pipeline(NearMiss(sampling_strategy='majority'), LOR_clf)
    us_model = us_pipeline.fit(us_xtrain[train], us_ytrain[train])
    us_predict= us_model.predict(us_xtrain[test])
for train, test in skf.split(us_xtrain, us_ytrain):
    us_pipeline =
    →imbalanced_make_pipeline(NearMiss(sampling_strategy='majority'),
    →decision_tree_clf)
    us_model = us_pipeline.fit(us_xtrain[train], us_ytrain[train])
    us_predict= us_model.predict(us_xtrain[test])
for train, test in skf.split(us_xtrain, us_ytrain):
    us_pipeline =
    →imbalanced_make_pipeline(NearMiss(sampling_strategy='majority'), KNN_clf)
    us_model = us_pipeline.fit(us_xtrain[train], us_ytrain[train])
    us_predict= us_model.predict(us_xtrain[test])
LR_predict = cross_val_predict(LOR_clf, train_x, train_y,
    →cv=5,method="decision_function")
KNC_predict = cross_val_predict(KNN_clf, train_x, train_y, cv=5)
DTC_predict = cross_val_predict(decision_tree_clf, train_x, train_y, cv=5)

LR_fpr, LR_tpr, LR_threshold = roc_curve(train_y, LR_predict)
KNN_fpr, KNN_tpr, KNN_threshold = roc_curve(train_y, KNC_predict)
DTC_fpr, DTC_tpr, DTC_threshold = roc_curve(train_y, DTC_predict)
plt.title('ROC Curve \n 3 Classifiers')
plt.plot(LR_fpr, LR_tpr, label='Logistic Regression Classifier' )
plt.plot(KNN_fpr, KNN_tpr, label='KNears Neighbors Classifier' )
plt.plot(DTC_fpr, DTC_tpr, label='Decision Tree Classifier' )
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')

```

```
plt.legend()
plt.show()
```



```
[24]: os_y_res = best_OS_LOR.decision_function(orig_x_test)

os_avg_precision = average_precision_score(orig_y_test, os_y_res)
#OS during Cross Validation
print('Average precision-recall of oversample by logestic regression(imprtant_
↪metric): {0:0.2f}'.format(
    os_avg_precision))

us_y_res = LOR_clf.decision_function(orig_x_test)

us_avg_precision = average_precision_score(orig_y_test, us_y_res)

US_predict= LOR_clf.predict(orig_x_test)
y_predict_LOR = LOR_clf.predict(test_x)
y_predict_KNN = KNN_clf.predict(test_x)
y_predict_DTC = decision_tree_clf.predict(test_x)
OS_predict = best_OS_LOR.predict(orig_x_test)
```

```

print('Average precision-recall of undersample by logistic regression(important_
↳metric): {0:0.2f}'.format(
    us_avg_precision))
print("misleading:")
print("MSE of undersample LOR: \n{}\n".format(mean_squared_error(test_y,
↳y_predict_LOR)))
print("MSE of undersample KNN: \n{}\n".format(mean_squared_error(test_y,
↳y_predict_KNN)))
print("MSE of undersample DTC: \n{}\n".format(mean_squared_error(test_y,
↳y_predict_DTC)))
print("MSE of oversample LOR of original data: \n{}\n".
↳format(mean_squared_error(orig_y_test, OS_predict)))
print("MSE of undersample LOR of original data: \n{}\n".
↳format(mean_squared_error(orig_y_test, US_predict)))

```

```

Average precision-recall of oversample by logistic regression(important metric):
0.77
Average precision-recall of undersample by logistic regression(important metric):
0.02
misleading:
MSE of undersample LOR:
0.17258883248730963

MSE of undersample KNN:
0.1065989847715736

MSE of undersample DTC:
0.4263959390862944

MSE of oversample LOR of original data:
0.020417478625726373

MSE of undersample LOR of original data:
0.25031161672021207

```

[]:

EE660 Project

December 3, 2020

```
[169]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb
import matplotlib.patches as mpatches
import collections
from collections import Counter
import warnings
warnings.filterwarnings("ignore")
```

```
[170]: from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import KFold, StratifiedKFold
```

```
[171]: data = pd.read_csv('creditcard.csv')
data
```

```
[171]:
```

	Time	V1	V2	V3	V4	V5	\		
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321			
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018			
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198			
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309			
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193			
...			
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473			
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229			
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515			
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961			
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546			
		V6	V7	V8	V9	...	V21	V22	\
0	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838		
1	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672		
2	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679		
3	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274		

4	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278
...
284802	-2.606837	-4.918215	7.305334	1.914428	...	0.213454	0.111864
284803	1.058415	0.024330	0.294869	0.584800	...	0.214205	0.924384
284804	3.031260	-0.296827	0.708417	0.432454	...	0.232045	0.578229
284805	0.623708	-0.686180	0.679145	0.392087	...	0.265245	0.800049
284806	-0.649617	1.577006	-0.414650	0.486180	...	0.261057	0.643078

	V23	V24	V25	V26	V27	V28	Amount	\
0	-0.110474	0.066928	0.128539	-0.189115	0.133558	-0.021053	149.62	
1	0.101288	-0.339846	0.167170	0.125895	-0.008983	0.014724	2.69	
2	0.909412	-0.689281	-0.327642	-0.139097	-0.055353	-0.059752	378.66	
3	-0.190321	-1.175575	0.647376	-0.221929	0.062723	0.061458	123.50	
4	-0.137458	0.141267	-0.206010	0.502292	0.219422	0.215153	69.99	
...
284802	1.014480	-0.509348	1.436807	0.250034	0.943651	0.823731	0.77	
284803	0.012463	-1.016226	-0.606624	-0.395255	0.068472	-0.053527	24.79	
284804	-0.037501	0.640134	0.265745	-0.087371	0.004455	-0.026561	67.88	
284805	-0.163298	0.123205	-0.569159	0.546668	0.108821	0.104533	10.00	
284806	0.376777	0.008797	-0.473649	-0.818267	-0.002415	0.013649	217.00	

	Class
0	0
1	0
2	0
3	0
4	0
...	...
284802	0
284803	0
284804	0
284805	0
284806	0

[284807 rows x 31 columns]

```
[173]: Total_value = len(data)
NF_value = data['Class'].value_counts()[0]
F_value = data['Class'].value_counts()[1]
print("No fraud number percentage", round(NF_value/Total_value*100,2))
print("fraud number percentage", round(F_value/Total_value*100,2))
```

No fraud number percentage 99.83
fraud number percentage 0.17

```
[174]: sb.countplot('Class', data=data)
plt.title('Class Distributions non fraud and fraud', fontsize=10)
```

```
[174]: Text(0.5, 1.0, 'Class Distributions non fraud and fraud')
```



```
[175]: from sklearn.preprocessing import RobustScaler
from sklearn.preprocessing import StandardScaler
rob = RobustScaler()
data['rs_amount'] = rob.fit_transform(data['Amount'].values.reshape(-1,1))
data['rs_time'] = rob.fit_transform(data['Time'].values.reshape(-1,1))
data.drop(['Time', 'Amount'],axis=1, inplace=True)

rs_amount = data['rs_amount']
rs_time = data['rs_time']

data.drop(['rs_amount', 'rs_time'],axis=1, inplace=True)
data.insert(0, 'rs_amount', rs_amount)
data.insert(1, 'rs_time', rs_time)

data
```

```
[175]:
```

	rs_amount	rs_time	V1	V2	V3	V4	\
0	1.783274	-0.994983	-1.359807	-0.072781	2.536347	1.378155	
1	-0.269825	-0.994983	1.191857	0.266151	0.166480	0.448154	
2	4.983721	-0.994972	-1.358354	-1.340163	1.773209	0.379780	
3	1.418291	-0.994972	-0.966272	-0.185226	1.792993	-0.863291	
4	0.670579	-0.994960	-1.158233	0.877737	1.548718	0.403034	

...
284802	-0.296653	1.034951	-11.881118	10.071785	-9.834783	-2.066656
284803	0.038986	1.034963	-0.732789	-0.055080	2.035030	-0.738589
284804	0.641096	1.034975	1.919565	-0.301254	-3.249640	-0.557828
284805	-0.167680	1.034975	-0.240440	0.530483	0.702510	0.689799
284806	2.724796	1.035022	-0.533413	-0.189733	0.703337	-0.506271

	V5	V6	V7	V8	...	V20	V21	\
0	-0.338321	0.462388	0.239599	0.098698	...	0.251412	-0.018307	
1	0.060018	-0.082361	-0.078803	0.085102	...	-0.069083	-0.225775	
2	-0.503198	1.800499	0.791461	0.247676	...	0.524980	0.247998	
3	-0.010309	1.247203	0.237609	0.377436	...	-0.208038	-0.108300	
4	-0.407193	0.095921	0.592941	-0.270533	...	0.408542	-0.009431	

...
284802	-5.364473	-2.606837	-4.918215	7.305334	...	1.475829	0.213454
284803	0.868229	1.058415	0.024330	0.294869	...	0.059616	0.214205
284804	2.630515	3.031260	-0.296827	0.708417	...	0.001396	0.232045
284805	-0.377961	0.623708	-0.686180	0.679145	...	0.127434	0.265245
284806	-0.012546	-0.649617	1.577006	-0.414650	...	0.382948	0.261057

	V22	V23	V24	V25	V26	V27	V28	\
0	0.277838	-0.110474	0.066928	0.128539	-0.189115	0.133558	-0.021053	
1	-0.638672	0.101288	-0.339846	0.167170	0.125895	-0.008983	0.014724	
2	0.771679	0.909412	-0.689281	-0.327642	-0.139097	-0.055353	-0.059752	
3	0.005274	-0.190321	-1.175575	0.647376	-0.221929	0.062723	0.061458	
4	0.798278	-0.137458	0.141267	-0.206010	0.502292	0.219422	0.215153	

...
284802	0.111864	1.014480	-0.509348	1.436807	0.250034	0.943651	0.823731
284803	0.924384	0.012463	-1.016226	-0.606624	-0.395255	0.068472	-0.053527
284804	0.578229	-0.037501	0.640134	0.265745	-0.087371	0.004455	-0.026561
284805	0.800049	-0.163298	0.123205	-0.569159	0.546668	0.108821	0.104533
284806	0.643078	0.376777	0.008797	-0.473649	-0.818267	-0.002415	0.013649

	Class
0	0
1	0
2	0
3	0
4	0
...	...
284802	0
284803	0
284804	0
284805	0
284806	0

[284807 rows x 31 columns]

```
[236]: #splitting data before undersample and oversample
from sklearn.model_selection import StratifiedShuffleSplit
from sklearn.model_selection import train_test_split
x = data.drop('Class', axis=1)
y = data['Class']
skf = StratifiedKFold(n_splits=5, random_state=None, shuffle=False)

for tr_index, te_index in skf.split(x, y):
    orig_x_train, orig_x_test = x.iloc[tr_index], x.iloc[te_index]
    orig_y_train, orig_y_test = y.iloc[tr_index], y.iloc[te_index]
orig_x_train = orig_x_train.values
orig_y_train = orig_y_train.values
orig_x_test = orig_x_test.values
orig_y_test = orig_y_test.values
```

```
[253]: data= data.sample(frac=1)
F_data = data.loc[data['Class'] == 1]
NF_data = data.loc[data['Class']==0][:492]
sub_data = pd.concat([F_data,NF_data]).sample(frac = 1, random_state = 42)

print('subsample distribution')
print(sub_data['Class'].value_counts()/len(sub_data))
sub_data
```

```
subsample distribution
1    0.5
0    0.5
Name: Class, dtype: float64
```

```
[253]:
```

	rs_amount	rs_time	V1	V2	V3	V4	V5	\
9262	-0.098232	-0.838297	1.226878	0.063952	0.795418	0.109193	-0.491986	
6427	-0.293440	-0.905579	0.725646	2.300894	-5.329976	4.007683	-1.730411	
133014	0.139733	-0.052656	1.135322	0.142540	0.122711	0.474670	-0.003014	
249963	-0.296653	0.821967	-0.679521	4.672553	-6.814798	7.143500	0.928654	
204079	1.208831	0.592230	1.862102	-0.124052	-1.989752	0.382609	0.473032	
...	
144754	4.216726	0.019784	-0.670238	0.945206	0.610051	2.640065	-2.707775	
247673	3.156012	0.810172	-5.192496	3.164721	-5.047679	2.246597	-4.011781	
61367	-0.279746	-0.409908	1.252850	0.293765	-0.170780	0.243902	0.528910	
151730	-0.042479	0.134435	-1.952933	3.541385	-1.310561	5.955664	-1.003993	
74794	4.051003	-0.339901	-6.003422	-3.930731	-0.007045	1.714669	3.414667	
	V6	V7	V8	...	V20	V21	V22	\
9262	-0.345963	-0.409757	-0.090610	...	-0.003457	-0.297851	-0.635229	
6427	-1.732193	-3.968593	1.063728	...	0.504646	0.589669	0.109541	
133014	-0.078456	-0.068025	0.101293	...	-0.035654	-0.213125	-0.688959	
249963	-1.873013	-2.306689	0.993702	...	0.872006	0.566849	-0.321691	

```


204079 -0.674517  0.298621 -0.282416 ...  0.150727 -0.204158 -0.511441
...
144754  1.952611 -1.624608 -5.229908 ...  1.474929 -2.504450  1.436472
247673 -0.638908 -2.873463  1.576318 ... -1.850470  1.167244 -1.006617
61367   0.413733 -0.052919  0.101599 ... -0.001047 -0.273833 -0.738917
151730  0.983049 -4.587235 -4.892184 ...  1.965030 -1.998091  1.133706
74794  -2.329583 -1.901512 -2.746111 ... -4.128186  1.101671 -0.992494

      V23      V24      V25      V26      V27      V28  Class
9262   0.094170 -0.000176  0.027330  0.748993 -0.089141 -0.001460      0
6427   0.601045 -0.364700 -1.843078  0.351909  0.594550  0.099372      1
133014 0.079722 -0.355197  0.143186  0.124495 -0.017226  0.018885      0
249963 -0.281325 -1.120256 -0.073394  0.553530  0.760542  0.386742      1
204079 0.077874  0.388335  0.007896 -0.120980 -0.019579  0.006155      1
...
144754 0.351542  0.648467  0.579681  0.075738  0.346717  0.282209      1
247673 0.774562  0.063397 -0.390658  1.884741 -1.742558 -0.082216      1
61367  -0.009814 -1.162879  0.301826  0.175835 -0.005637  0.007501      0
151730 -0.041461 -0.215379 -0.865599  0.212545  0.532897  0.357892      1
74794  -0.698259  0.139898 -0.205151 -0.472412  1.775378 -0.104285      1

```

[984 rows x 31 columns]

```
[254]: sb.heatmap(sub_data.corr(), cmap='coolwarm_r', annot_kws={'size':20})
plt.show()
```



```
[255]: from sklearn.model_selection import train_test_split
sub_y = sub_data['Class']
sub_x = sub_data.drop('Class', axis = 1)
train_x, test_x, train_y, test_y = train_test_split(sub_x, sub_y, test_size=0.
↳2, random_state=42)
train_x = train_x.values
train_y = train_y.values
test_x = test_x.values
test_y = test_y.values
```

```
[274]: #baseline 1
from sklearn.model_selection import cross_val_score
DTC = DecisionTreeClassifier()
DTC.fit(train_x, train_y)
train_score = cross_val_score(DTC, train_x, train_y, cv=5)
avg_score = train_score.mean()
print("decision tree Classifier training score ", round(avg_score, 2)*100)
```

decision tree Classifier training score 89.0

```
[275]: from sklearn.model_selection import GridSearchCV
#thanks to internet source, i got valid parameter choices
para_tree = {"criterion": ["gini", "entropy"], "max_depth": list(range(2,4,1)),
             "min_samples_leaf": list(range(5,7,1))}
decision_tree = GridSearchCV(DecisionTreeClassifier(), para_tree)
decision_tree.fit(train_x, train_y)
decision_tree_clf = decision_tree.best_estimator_
decision_tree_score = cross_val_score(decision_tree_clf, train_x, train_y, cv=5)
avg_tree_score = decision_tree_score.mean()
print("decision tree Classifier cross validation score",
↳round(avg_tree_score*100, 2).astype(str))
```

decision tree Classifier cross validation score 91.87

```
[276]: from sklearn.pipeline import make_pipeline
from imblearn.pipeline import make_pipeline as imbalanced_make_pipeline
from imblearn.over_sampling import SMOTE
from imblearn.under_sampling import NearMiss
from imblearn.metrics import classification_report_imbalanced
us_x = data.drop('Class', axis=1)
us_y = data['Class']
# the following make pipelined codes was modified from www.kaggle.com/
↳janiobachmann
for tr_index, te_index in skf.split(us_x, us_y):
    us_xtrain, us_xtest = us_x.iloc[tr_index], us_x.iloc[te_index]
```

```

us_ytrain, us_ytest = us_y.iloc[tr_index], us_y.iloc[te_index]


us_xtrain = us_xtrain.values
us_xtest = us_xtest.values
us_ytrain = us_ytrain.values
us_ytest = us_ytest.values
# the following make pipelined codes was modified from www.kaggle.com/
→janiobachmann
for train, test in skf.split(us_xtrain, us_ytrain):
    us_pipeline =
    →imbalanced_make_pipeline(NearMiss(sampling_strategy='majority'),
    →decision_tree_clf)
    us_model = us_pipeline.fit(us_xtrain[train], us_ytrain[train])
    us_predict= us_model.predict(us_xtrain[test])

```

```

[277]: from sklearn.model_selection import ShuffleSplit
from sklearn.model_selection import learning_curve
cv = ShuffleSplit(n_splits=100, test_size=0.2, random_state=233)
tr_sizes=np.linspace(.1, 1.0, 5)
tr_sizes, tr_scores, te_scores = learning_curve(
    decision_tree_clf , train_x, train_y, cv=cv, n_jobs=1,
    →train_sizes=tr_sizes)
tr_scores_avg = np.mean(tr_scores,axis = 1)
te_scores_avg = np.mean(te_scores,axis = 1)
plt.plot(tr_sizes, tr_scores_avg, 'o-', color="blue",
        label="Training score")
plt.plot(tr_sizes, te_scores_avg, 'x-', color="red",
        label="Cross-validation score")
plt.title("DTC Learning Curve")
plt.xlabel('Training size')
plt.ylabel('Score')
plt.grid(True)
plt.legend(loc="best")
plt.show()

```



```
[278]: from sklearn.metrics import roc_curve
from sklearn.model_selection import cross_val_predict
from sklearn.metrics import roc_auc_score
DTC_predict = cross_val_predict(decision_tree_clf, train_x, train_y, cv=5)
print('DTC: ', roc_auc_score(train_y, DTC_predict))
```

DTC: 0.916631116282076

```
[318]: print("Best Parameters of DTC: \n{}\n".format(decision_tree.best_params_))
```

Best Parameters of DTC:
{'criterion': 'gini', 'max_depth': 3, 'min_samples_leaf': 5}

```
[283]: from sklearn.metrics import recall_score
from sklearn.metrics import precision_score
from sklearn.metrics import f1_score
from sklearn.metrics import accuracy_score

y_predict_DTC = decision_tree_clf.predict(test_x)

print('Recall of DTC : {:.2f}'.format(recall_score(test_y, y_predict_DTC)))
print('Precision of DTC : {:.2f}'.format(precision_score(test_y,
↪y_predict_DTC)))
```

```
print('F1 of DTC: {:.2f}'.format(f1_score(test_y,y_predict_DTC)))
print('Accuracy of DTC : {:.2f}'.format(accuracy_score(test_y, y_predict_DTC)))
```

Recall of DTC : 0.97
Precision of DTC : 0.66
F1 of DTC: 0.79
Accuracy of DTC : 0.71

```
[284]: #baseline2
KNC = KNeighborsClassifier()
KNC.fit(train_x, train_y)
train_score_KNC = cross_val_score(KNC, train_x, train_y, cv=5)
avg_score_KNC = train_score.mean()
print("KNN Classifier training score ", round(avg_score_KNC*100, 2))
#thanks to internet source, i got valid parameter choices
para_KNC = {"n_neighbors": list(range(2,5,1)), 'algorithm': ['auto',
    ↳'ball_tree', 'kd_tree', 'brute']}
KNN = GridSearchCV(KNeighborsClassifier(), para_KNC)
KNN.fit(train_x, train_y)
KNN_clf = KNN.best_estimator_
KNN_score = cross_val_score(KNN_clf, train_x, train_y, cv=5)
avg_KNN_score = KNN_score.mean()
print("KNN Classifier cross validation score", round(avg_KNN_score*100, 2).
    ↳astype(str))
```

KNN Classifier training score 88.69
KNN Classifier cross validation score 92.88


```
[285]: # the following make pipelined codes was modified from www.kaggle.com/
    ↳janiobachmann
for train, test in skf.split(us_xtrain, us_ytrain):
    us_pipeline =
    ↳imbalanced_make_pipeline(NearMiss(sampling_strategy='majority'), KNN_clf)
    us_model = us_pipeline.fit(us_xtrain[train], us_ytrain[train])
    us_predict= us_model.predict(us_xtrain[test])
```

```
[286]: tr_sizes_KNN=np.linspace(.1, 1.0, 5)
tr_sizes_KNN, tr_scores_KNN, te_scores_KNN = learning_curve(
    KNN_clf , train_x, train_y, cv=cv, n_jobs=1, train_sizes=tr_sizes_KNN)
tr_scores_KNN_avg = np.mean(tr_scores_KNN,axis = 1)
te_scores_KNN_avg = np.mean(te_scores_KNN,axis = 1)
plt.plot(tr_sizes, tr_scores_KNN_avg, 'o-', color="blue",
    label="Training score")
plt.plot(tr_sizes, te_scores_KNN_avg, 'x-', color="red",
    label="Cross-validation score")
plt.title("KNC Learning Curve")
plt.xlabel('Training size')
```

```

plt.ylabel('Score')
plt.grid(True)
plt.legend(loc="best")
plt.show()
KNC_predict = cross_val_predict(KNN_clf, train_x, train_y, cv=5)
print('KNNC: ', roc_auc_score(train_y, KNC_predict))

```



KNNC: 0.927370564281559

```

[290]: y_predict_KNN = KNN_clf.predict(test_x)

print('Recall of KNN : {:.2f}'.format(recall_score(test_y, y_predict_KNN)))
print('Precision of KNN : {:.2f}'.format(precision_score(test_y, y_predict_KNN)))
print('F1 of KNN: {:.2f}'.format(f1_score(test_y, y_predict_KNN)))
print('Accuracy of KNN : {:.2f}'.format(accuracy_score(test_y, y_predict_KNN)))

```

Recall of KNN : 0.95
Precision of KNN : 0.91
F1 of KNN: 0.93
Accuracy of KNN : 0.92

```

[325]: print("Best Parameters of KNN: \n{}\n".format(KNN.best_params_))

```


Best Parameters of KNN:

```
{'algorithm': 'auto', 'n_neighbors': 4}
```

```
[292]: # good one
LR = LogisticRegression()
LR.fit(train_x, train_y)
train_score_LR = cross_val_score(LR, train_x, train_y, cv=5)
avg_score_LR = train_score_LR.mean()
print("Logistic regression training score ", round(avg_score_LR*100, 2))
#thanks to internet source, i got valid parameter choices
para_LR = {"penalty": ['l1', 'l2'], 'C': [0.001, 0.01, 0.1, 1, 10, 100, 1000]}
LOR = GridSearchCV(LogisticRegression(), para_LR)
LOR.fit(train_x, train_y)
LOR_clf = LOR.best_estimator_
LOR_score = cross_val_score(LOR_clf, train_x, train_y, cv=5)
avg_LOR_score = LOR_score.mean()
print("Logistic regression cross validation score", round(avg_LOR_score*100, 2)).
      ↳astype(str))
```

Logistic regression training score 88.69

Logistic regression cross validation score 93.77


```
[293]: from sklearn.pipeline import make_pipeline
from imblearn.pipeline import make_pipeline as imbalanced_make_pipeline
from imblearn.over_sampling import SMOTE
from imblearn.under_sampling import NearMiss
from imblearn.metrics import classification_report_imbalanced
# the following make pipelined codes was modified from www.kaggle.com/
  ↳janiobachmann
for train, test in skf.split(us_xtrain, us_ytrain):
    us_pipeline =
    ↳imbalanced_make_pipeline(NearMiss(sampling_strategy='majority'), LOR_clf)
    us_model = us_pipeline.fit(us_xtrain[train], us_ytrain[train])
    us_predict= us_model.predict(us_xtrain[test])
```

```
[294]: tr_sizes_LR=np.linspace(.1, 1.0, 5)
tr_sizes_LR, tr_scores_LR, te_scores_LR = learning_curve(
    LOR_clf , train_x, train_y, cv=cv, n_jobs=1, train_sizes=tr_sizes_LR)
tr_scores_LR_avg = np.mean(tr_scores_LR,axis = 1)
te_scores_LR_avg = np.mean(te_scores_LR,axis = 1)
plt.plot(tr_sizes_LR, tr_scores_LR_avg, 'o-', color="blue",
        label="Training score")
plt.plot(tr_sizes_LR, te_scores_LR_avg, 'x-', color="red",
        label="Cross-validation score")
plt.title("LR Learning Curve")
plt.xlabel('Training size')
```

```

plt.ylabel('Score')
plt.grid(True)
plt.legend(loc="best")
plt.show()
LR_predict = cross_val_predict(LOR_clf, train_x, train_y, u
    ↪cv=5,method="decision_function")
print('LORC: ', roc_auc_score(train_y, LR_predict))

```



LORC: 0.9726585223967423

```

[295]: y_predict_LOR = LOR_clf.predict(test_x)

print('Recall of LR : {:.2f}'.format(recall_score(test_y, y_predict_LOR)))
print('Precision of LR : {:.2f}'.format(precision_score(test_y, y_predict_LOR)))
print('F1 of LR: {:.2f}'.format(f1_score(test_y, y_predict_LOR)))
print('Accuracy of LR : {:.2f}'.format(accuracy_score(test_y, y_predict_LOR)))

```

Recall of LR : 0.98
Precision of LR : 0.69
F1 of LR: 0.81
Accuracy of LR : 0.74

```
[296]: from sklearn.metrics import average_precision_score
us_y_res = LOR_clf.decision_function(orig_x_test)

us_avg_precision = average_precision_score(orig_y_test, us_y_res)

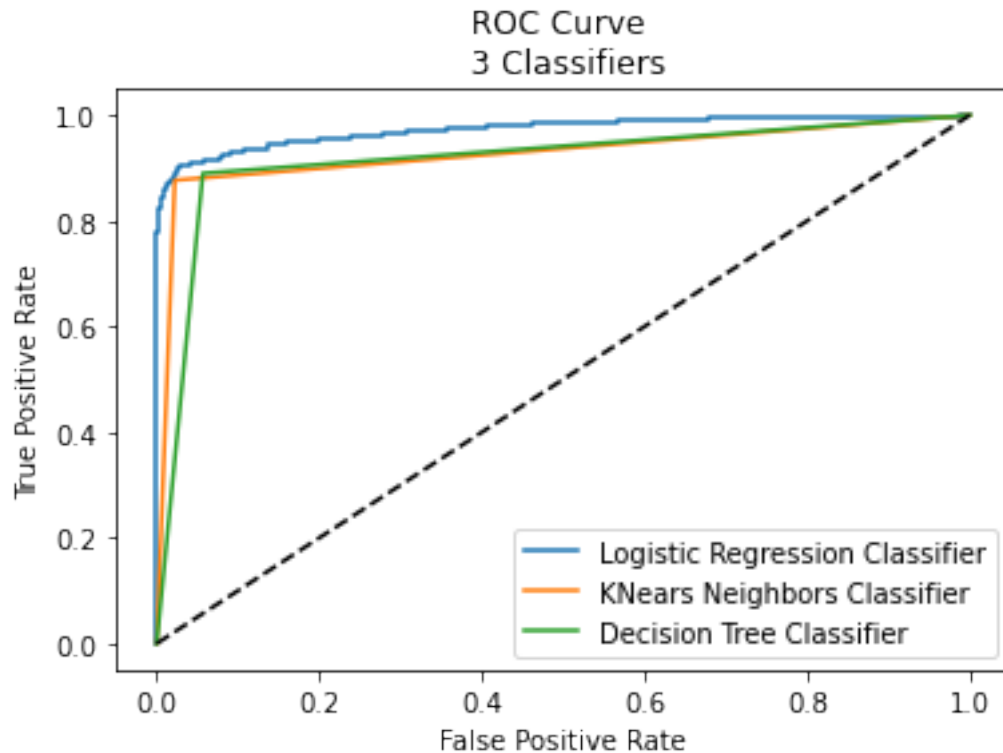
print('Average precision-recall: {0:0.2f}'.format(
    us_avg_precision))
```

Average precision-recall: 0.08

```
[323]: print("Best Parameters of under sample LR: \n{}\n".format(LOR.best_params_))
```

Best Parameters of under sample LR:
{'C': 1, 'penalty': 'l2'}

```
[333]: from sklearn.metrics import roc_curve
LR_fpr, LR_tpr, LR_threshold = roc_curve(train_y, LR_predict)
KNN_fpr, KNN_tpr, KNN_threshold = roc_curve(train_y, KNC_predict)
DTC_fpr, DTC_tpr, DTC_threshold = roc_curve(train_y, DTC_predict)
plt.title('ROC Curve \n 3 Classifiers')
plt.plot(LR_fpr, LR_tpr, label='Logistic Regression Classifier' )
plt.plot(KNN_fpr, KNN_tpr, label='KNears Neighbors Classifier' )
plt.plot(DTC_fpr, DTC_tpr, label='Decision Tree Classifier' )
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()
```



```
[335]: from imblearn.over_sampling import SMOTE
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import classification_report
LOR_sm = LogisticRegression()
rand_LOR = RandomizedSearchCV(LogisticRegression(), para_LR, n_iter=4)
# the following make pipelined codes was modified from www.kaggle.com/
# janiobachmann
for train, test in skf.split(orig_x_train, orig_y_train):
    os_pipeline = imbalanced_make_pipeline(SMOTE(sampling_strategy='minority'),
    rand_LOR)
    os_model = os_pipeline.fit(orig_x_train[train], orig_y_train[train])
    best_OS_LOR = rand_LOR.best_estimator_

OS_predict = best_OS_LOR.predict(orig_x_test)
```

```
[304]: os_y_res = best_OS_LOR.decision_function(orig_x_test)

os_avg_precision = average_precision_score(orig_y_test, os_y_res)
#OS during Cross Validation
print('Average precision-recall of oversample by logestic regression: {0:0.2f}'.
    format(
        os_avg_precision))
```

Average precision-recall of oversample by logistic regression: 0.73

```
[324]: print("Best Parameters of over sample: \n{}\n".format(rand_LOR.best_params_))
```

Best Parameters of over sample:

{'penalty': 'l2', 'C': 0.001}

```
[308]: #OS after Cross Validation
from imblearn.over_sampling import SMOTE
from sklearn.metrics import confusion_matrix
sm_os = SMOTE(sampling_strategy='minority', random_state=42)
sm_train_x, sm_train_y = sm_os.fit_sample(orig_x_train, orig_y_train)
sm_os_LOR = LOR.best_estimator_
sm_os_LOR.fit(sm_train_x, sm_train_y)
LOR_y_predict = sm_os_LOR.predict(test_x)
confusionM_of_LOR_sm_after_CV = confusion_matrix(test_y, LOR_y_predict)
confusionM_of_KNN = confusion_matrix(test_y, y_predict_KNN)
confusionM_of_DTC = confusion_matrix(test_y, y_predict_DTC)
```

```
[334]: from sklearn.metrics import mean_squared_error
US_predict= LOR_clf.predict(orig_x_test)
print("MSE of unersample LOR: \n{}\n".format(mean_squared_error(test_y,
    ↳y_predict_LOR)))
print("MSE of unersample KNN: \n{}\n".format(mean_squared_error(test_y,
    ↳y_predict_KNN)))
print("MSE of unersample DTC: \n{}\n".format(mean_squared_error(test_y,
    ↳y_predict_DTC)))
print("MSE of oversample LOR of original data: \n{}\n".
    ↳format(mean_squared_error(orig_y_test, OS_predict)))
print("MSE of undersample LOR of original data: \n{}\n".
    ↳format(mean_squared_error(orig_y_test, US_predict)))
```

MSE of unersample LOR:

0.25888324873096447

MSE of unersample KNN:

0.07614213197969544

MSE of unersample DTC:

0.29441624365482233


MSE of oversample LOR of original data:

0.024630887800424852

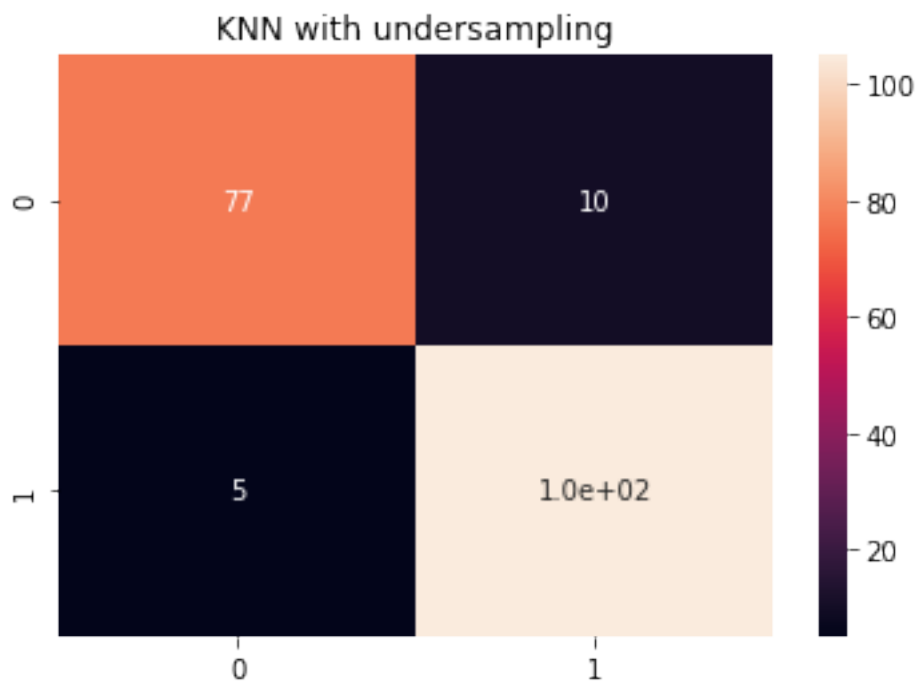
MSE of undersample LOR of original data:

0.024683555415108582

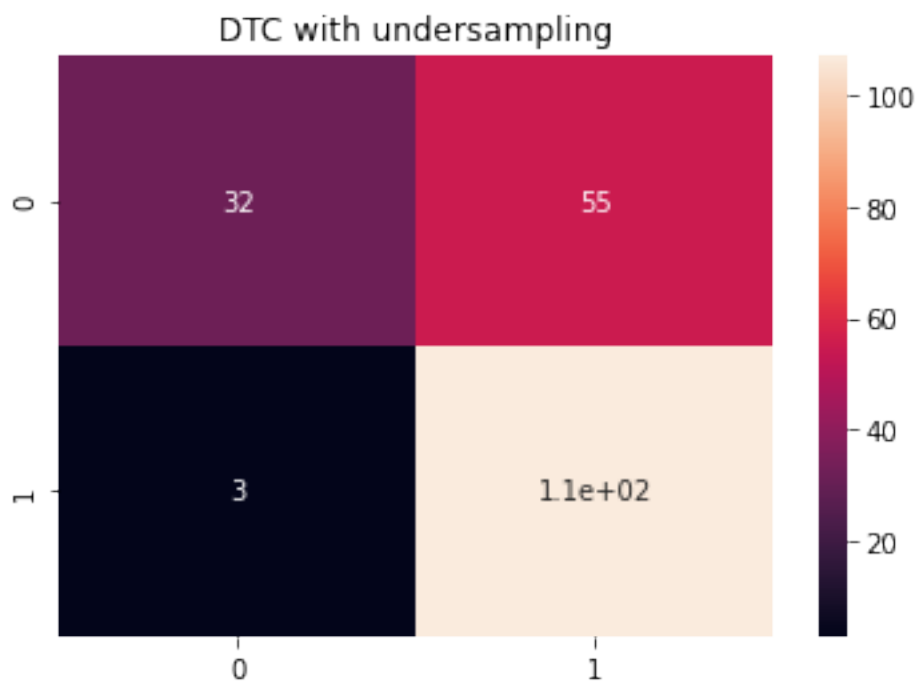
```
[312]: sb.heatmap(confusionM_of_LOR_sm_after_CV,annot=True)
plt.title("LR with oversampling after CV in subsample result")
plt.show()
```



```
[313]: sb.heatmap(confusionM_of_KNN,annot=True)
plt.title("KNN with undersampling in subsample result")
plt.show()
```



```
[316]: sb.heatmap(confusionM_of_DTC,annot=True)  
plt.title("DTC with undersampling in subsample result")  
plt.show()
```



```
[337]: import pickle
Pkl_Filename_LR_U = "Pickle_LR_undersample.pkl"
with open(Pkl_Filename_LR_U, 'wb') as file:
    pickle.dump(LOR_clf, file)
Pkl_Filename_LR_O = "Pickle_LR_oversample.pkl"
with open(Pkl_Filename_LR_O, 'wb') as file:
    pickle.dump(best_OS_LOR, file)
Pkl_Filename_KNN = "Pickle_KNN_undersample.pkl"
with open(Pkl_Filename_KNN, 'wb') as file:
    pickle.dump(KNN_clf, file)
Pkl_Filename_DTC = "Pickle_DTC_undersample.pkl"
with open(Pkl_Filename_DTC, 'wb') as file:
    pickle.dump(decision_tree_clf, file)
```

```
[353]: #The folloing plotting code changed from
#https://scikit-learn.org/stable/auto_examples/ensemble/
→plot_voting_decision_regions.html
dbx=sub_x[["V12", "V14"]]
dby=sub_y
LOR_clf.fit(dbx,dby)
KNN_clf.fit(dbx,dby)
decision_tree_clf.fit(dbx,dby)


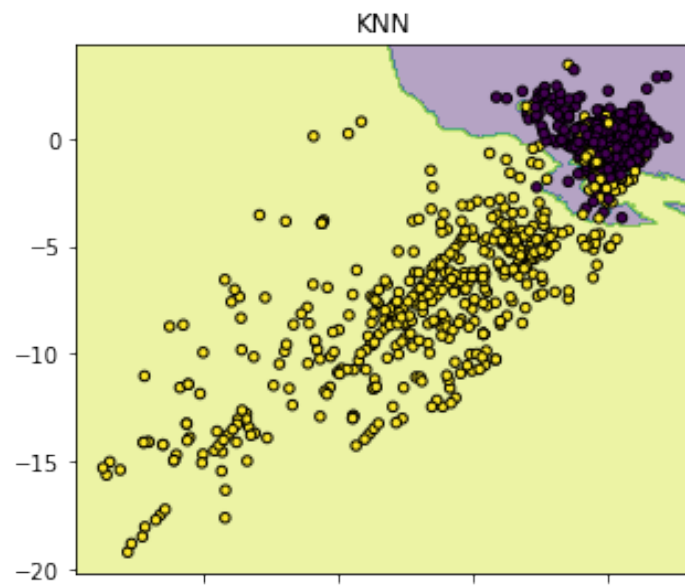

x_min= dbx.to_numpy()[:,0].min() - 1
x_max= dbx.to_numpy()[:,0].max() + 1
y_min, y_max = dbx.to_numpy()[:, 1].min() - 1, dbx.to_numpy()[:, 1].max() + 1
x_x, y_y = np.meshgrid(np.arange(x_min, x_max, 0.1),
                        np.arange(y_min, y_max, 0.1))

f, axarr = plt.subplots(3, 1, sharex='col', sharey='row',figsize=(5, 15))

for index, classifier, title in zip([0,1,2],
                                    [LOR_clf, KNN_clf, decision_tree_clf],
                                    ['Logistic Regression', 'KNN',
                                    'Decision tree']):

    Z = classifier.predict(np.c_[x_x.ravel(), y_y.ravel()])
    Z = Z.reshape(x_x.shape)

    axarr[index].contourf(x_x, y_y, Z, alpha=0.4)
    axarr[index].scatter(dbx.to_numpy()[:, 0], dbx.to_numpy()[:, 1], c=dby.
        →to_numpy(),
                        s=20, edgecolor='k')
    axarr[index].set_title(title)
plt.show()
```

[]: